

Hub-and-Spoke — LLM-as-Judge Report

Council: Hub-and-Spoke (Judge asks agents in discussion rounds) · Judge: Qwen/Qwen3.6-35B-A3B-FP8 · Dataset: 500 images

EXECUTIVE SUMMARY

64.6%

COUNTRY ACCURACY

435 km

MEDIAN HAVERSINE

1598 km

MEAN HAVERSINE

500

IMAGES

Geo-spatial Bias

N/S


strong north bias (p=0.0005) (p=0.0005)

E/W

no significant bias (p=0.8940) (p=0.8940)

Quadrants


SW: 116 · NE: 142 · NW: 135 · SE: 107




Top Confusion Pairs

TRUTH	PREDICTED	COUNT
canada	united states	5
south africa	botswana	4
south africa	namibia	3
australia	new zealand	3
peru	colombia	3
uruguay	brazil	3
indonesia	philippines	2
argentina	mexico	2

TRUTH	PREDICTED	COUNT
czechia	poland	2
poland	lithuania	2



Per-Agent Metrics



LLM-as-Judge Scores

AGENT	N	ROLE	HALLUC.↓	VIS.CONS.↑	CALIB.↑
linguistic	500	100.0%	0.01	0.99	0.58
landscape	500	100.0%	0.11	0.90	0.65
botanics	500	100.0%	0.09	0.89	0.64
regulatory	500	100.0%	0.09	0.92	0.63
meta	500	100.0%	0.15	0.85	0.62

